

# Chance, Machine Learning Algorithms, MVPA, Brains on Trial

Commentary on this assignment:

This assignment revisits concepts from probability and hypothesis testing in assessing accuracy of machine learning algorithms. Then you'll investigate a specific ML algorithm (separate submission). Then you'll apply machine learning to a more complex problem than the 4 brain regions, which will allow you to apply methods to avoid overfitting. Finally, you'll read or watch a video about the use of neuroscience in the courtroom and in lie detection, which raises important questions about how our legal system works.

I provided very rough time estimates just to give you a sense of scale. I expect that there will be lots of variability in how much time people spend on this assignment overall and how much time people spend in each section. I trust you to make the good decisions for your learning goals and the course goals around challenging yourself, engaging your curiosity, and contributing to class.

This document is a bit long, so I've included a linked table of contents and pictures of puppies (puppies only available in this version: <https://docs.google.com/document/d/1CmVUOSBBxdDdp-JU8pCzN5pKm6gdCE1sWJtiQXHqv4E/edit?usp=sharing>).

The data you need are here: <https://drive.google.com/file/d/12cJmNYllwkhTGyHa9NAWpZWJVqPw4hib/view?usp=sharing>

## Table of Contents

Chance, Machine Learning Algorithms, MVPA, Brains on Trial.....	1
Understanding "chance" in relation to classification [30 mins].....	2
Exercise 1a: What is the most probable accuracy?.....	2
Exercise 1b: What is the expected value of the accuracy?.....	3
Exercise 2a: For the training data, what is the probability of getting an accuracy of 50% or worse?.....	3
Exercise 2b: For the training data, what is the probability of getting an accuracy of 50% or worse?.....	3
Exercise 3a: What should I choose as my accuracy threshold for the training data? .....	3
Exercise 3b: What accuracy threshold should I use for my test data?.....	4
Exercise 3c: Are these the same or different? Why?.....	4
Exercise 4: What is the probability that it will perform at chance or worse on the test data?.....	4
Exercise 5: What classification accuracy would we get on the train data if our algorithm just guessed face every time?.....	4
Challenge exercise: If our algorithm is still just a coin flipper, what is the expected value of the accuracy on our training data?.....	4
Use external resources to investigate, understand, and explain an algorithm [90 mins].....	5
Exercise 6: Explain algorithm (separate submission).....	5
Multi-voxel pattern analysis on fMRI data [90 mins].....	5
Exercise (BIG ONE): Run multivoxel pattern analysis (machine learning on voxels).....	7
Brains on Trial [120 minutes].....	15

## Understanding “chance” in relation to classification [30 mins]

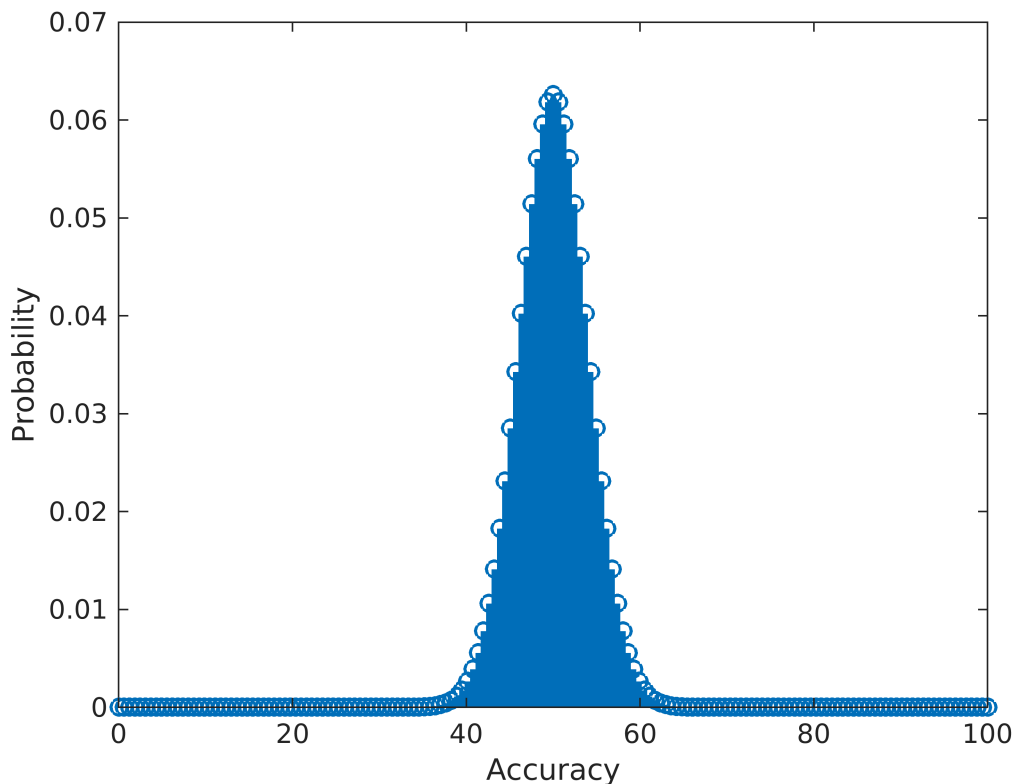
In this section, we’re going to explore probability in relation to classification accuracy. We’ll use the example from our faces vs houses fMRI data. In the training data, we have 162 total examples and in the test data we have 54 total rows of data. Each is evenly split between faces and houses.

If my algorithm just flips a coin each time (“operates at chance level”,  $p(\text{face})=0.5$ ):

### Exercise 1a: What is the most probable accuracy?

I think it should be close to 50, but since we have a discrete sample set, the accuracy might not be exactly 50 if there is an odd number of samples

```
n = 162;
test_n = 54;
b = makedist("Binomial", 'N', n, 'p', 0.5);
x = 0:n; % number of correct predictions
stem(100.*(x./n), b.pdf(x), 'LineWidth', 2);
xlabel('Accuracy');
ylabel('Probability');
```



## Exercise 1b: What is the expected value of the accuracy?

Here are some resources about expected values that might be helpful:

[https://ocw.mit.edu/courses/mathematics/18-05-introduction-to-probability-and-statistics-spring-2014/readings/MIT18\\_05S14\\_Reading4b.pdf](https://ocw.mit.edu/courses/mathematics/18-05-introduction-to-probability-and-statistics-spring-2014/readings/MIT18_05S14_Reading4b.pdf)

<https://www.stat.purdue.edu/~zhanghao/STAT511/handout/Stt511%20Sec3.3.pdf>

Expecting around 50, but it can be a non-discrete integer because it is a weighted average

```
expectedValueAccuracy = b.mean/n
```

```
expectedValueAccuracy = 0.5000
```

## Exercise 2a: For the training data, what is the probability of getting an accuracy of 50% or worse?

```
p = binocdf(n*0.5, n, 0.5) % inclusive of 50%
```

```
p = 0.5313
```

## Exercise 2b: For the training data, what is the probability of getting an accuracy of 50% or worse?

```
% for testing data
```

```
p = binocdf(test_n*0.5, test_n, 0.5) % inclusive of 50% and fewer options so each option
```

```
p = 0.5540
```

I want to set a threshold to reject the hypothesis that my algorithm is operating a chance level. Let's say that I feel comfortable being wrong 10% of the time when my algorithm is operating at chance.

## Exercise 3a: What should I choose as my accuracy threshold for the training data?

null\_hypothesis: my algorithm is operating at chance level

alternative\_hypothesis: my algorithm is not operating at chance level

alpha = 0.1

```
b = makedist("Binomial", 'N', n, 'p', 0.5);  
100 * icdf(b, 0.9)/n % if 90% were correct, find number correct
```

```
ans = 54.9383
```

### Exercise 3b: What accuracy threshold should I use for my test data?

```
b = makedist("Binomial", 'N', test_n, 'p', 0.5);  
100 * icdf(b, 0.9)/test_n % if 90% were correct, find number correct
```

```
ans = 59.2593
```

### Exercise 3c: Are these the same or different? Why?

Different because there are a different number of trials in train and test. If there are more trials, we can be more confident in our result and tighten the range.

Now, let's suppose I have an algorithm that is 70% accurate in reality.

### Exercise 4: What is the probability that it will perform at chance or worse on the test data?

```
p = 0.7;  
binocdf(0.5*test_n, test_n, p)
```

```
ans = 0.0016
```

In the questions above, we have been considering a situation where we have an equal number of face and house examples in our data. Now let's think about a situation where that is not the case. Our training data still has 162 total examples, but now 130 of them are faces.

### Exercise 5: What classification accuracy would we get on the train data if our algorithm just guessed face every time?

```
100 * 130/162
```

```
ans = 80.2469
```

### Challenge exercise: If our algorithm is still just a coin flipper, what is the expected value of the accuracy on our training data?

```
p = 0.5;  
tp = 130 * p; % number of tp guessed correctly by random chance  
tn = (162 - 130) * p; % number of tn guessed correctly by random chance  
(tp + tn)./162
```

```
ans = 0.5000
```

Used a bayesian tree to draw prior and posterior

.3/ \.7

.5/\.5 .5/\.5

## Use external resources to investigate, understand, and explain an algorithm [90 mins]

### Exercise 6: Explain algorithm (separate submission)

Choose one machine learning algorithm that you want to learn more about. Create a short explanation of this algorithm using whatever media you prefer (text, video, slides, whiteboard, audio, stone carving).

Please share any good resources that you find in this mutual document.

<https://docs.google.com/document/d/1AC4ULXd8sHEbw2M3u-tdkT0lvJyINvOjN7ojFg1NOQw/edit?usp=sharing>

If you are new to machine learning, I recommend “K nearest neighbors” or “linear support vector machines” as good starting algorithms.

**How to submit this part:** This part of the assignment has a separate submission. This will serve as one of our learning check-offs. Please submit this on time if possible, since we’ll probably be doing a peer review between Monday and Wednesday.

## Multi-voxel pattern analysis on fMRI data [90 mins]

This is the fMRI data from one person (subject 2 in this study: [Haxby, J.V., Gobbini, M.I., Furey, M.L., Ishai, A., Schouten, J.L., Pietrini, P. \(2001\). Distributed and overlapping representations of faces and objects in ventral temporal cortex. Science, 293\(5539\):2425-30](#)). There are 162 examples in the training set (these correspond to times where a person saw a house or a face). The labels tell you whether they were looking at a face or a house (same as last assignment). The data are linked at the top of this document.

Think of the voxel values as the potential features. In the train\_fmri3d data, the voxels are in a 3 dimensional array (that is 40 x 64 x 64), plus a 4th dimension that is the 162 examples. However, for the classifier, you need to reshape the data to a 2 dimensional array.

```
fmri2d= squeeze(reshape(train_fmri3d,1,1,[],162))'
```

```
fmri2d = 162x163840 single matrix
```

```
0    0    0    0    0    0    0    0    0    0    0    0    0 ...
0    0    0    0    0    0    0    0    0    0    0    0    0
0    0    0    0    0    0    0    0    0    0    0    0    0
0    0    0    0    0    0    0    0    0    0    0    0    0
0    0    0    0    0    0    0    0    0    0    0    0    0
0    0    0    0    0    0    0    0    0    0    0    0    0
0    0    0    0    0    0    0    0    0    0    0    0    0
0    0    0    0    0    0    0    0    0    0    0    0    0
0    0    0    0    0    0    0    0    0    0    0    0    0
```

⋮

fmri2d has the dimensions of examples x voxels. However, we know that a bunch of these values are outside of the brain (since train\_fmri3d is just a cube shape, and fun fact... our brains are not cube shaped).

If we wanted to make a variable that contained every voxel in the brain (but excluded those not in the brain), we can use our brain mask. The brain mask has 1s for voxels in the brain and 0s for voxels not in the brain. We need to reshape this to align with the shape of fmri2d.

```
brainmask2d=single(reshape(mask,1,[]))
```

```
brainmask2d = 1x163840 single row vector
```

```
0 0 0 0 0 0 0 0 0 0 0 0 0 ...
```

Create a variable called brainvals that has the dimensions of (examples x voxels in the brain). Exercise 7: What are the actual dimensions of this variable (numbers)?

```
brainvals = fmri2d.*repmat(brainmask2d, [162 1])
```

```
brainvals = 162x163840 single matrix
```

```
0 0 0 0 0 0 0 0 0 0 0 0 0 ...
0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0
⋮
```

```
brainvals = brainvals( :, any(brainvals))
```

```
brainvals = 162x39912 single matrix
```

```
45.1284 -26.6410 -42.3341 -15.2561 59.0485 -11.6695 -17.1604 -20.3824 ...
41.1242 -16.6047 -36.3399 -8.2558 63.0427 -18.6882 -23.0782 -6.3803
39.1200 -22.5684 -26.3456 -1.2556 62.0368 -20.7068 -20.9959 -12.3783
39.1159 -15.5321 -31.3513 -10.2553 68.0310 -15.7255 -26.9137 -17.3762
36.1117 -24.4957 -27.3571 -4.2551 75.0251 -7.7441 -22.8314 -15.3742
48.1075 -27.4594 -32.3628 -6.2549 77.0193 -10.7627 -21.7491 -19.3722
42.1034 -24.4231 -22.3686 0.7454 65.0134 -16.7814 -20.6669 -14.3701
45.0992 -17.3868 -27.3743 -10.2544 68.0076 -15.8000 -23.5846 -11.3681
40.0950 -21.3505 -19.3801 -10.2542 64.0017 -14.8186 -17.5024 -18.3660
31.9532 -48.1158 -42.5753 -56.2460 36.8028 -23.4524 -26.7056 -42.2967
⋮
```

162 x 39912

Dang, that's a lot of voxels! We could throw all of these voxels into our classifier, with each voxel serving as a feature. I would strongly recommend against this, why?

Think. Then highlight text below for answer

This may lead to overfitting our classifier. We have 162 examples and more than 30k variables.

Also, this will take forever to run and load.

## **Exercise (BIG ONE): Run multivoxel pattern analysis (machine learning on voxels).**

Run classification of faces vs houses using some subset of voxels. Be sure to use cross-validation to help prevent overfitting.

For this part of the assignment, I'm asking you to piece together and extend some of the code that you've seen already. You can look at the `brainViewer.mlx` for more example code for selecting a subset of voxels from the data. You can look at `fmriToFeaturesClassification` for tips on how to implement the classifier using only code (though you'll still need to look up how to do cross validation). You are also welcome to use the `classificationLearner` app if you prefer.

You should explore different aspects of the feature space and document what you try and what you observed. Some things you might try:

- downsample the brain data to a smaller number of voxels randomly
- downsample the brain data to a smaller number of voxels by strategically selecting a brain region (you may still want to downsample within a region if it is big)
- use a dimensionality reduction technique such as PCA or ICA
- create new features that represent and summarize information about groups of voxels (this is like what we did by taking the average activity in each of the 4 regions, but you can get more creative if you want to)

Challenge: Choose some parameter to manipulate (e.g., # of nearest neighbors in KNN algorithm, or how many voxels to include). Generate a plot to represent the relationship between this parameter and a measure of how well your classifier performs (e.g., accuracy). Important tips (feel free to add some here or ask questions)

- As I mentioned in class, the ventral temporal lobe plays an important role in visual processing. One of the 4 regions that we explored before is the ventral temporal lobe.
- If you are using the Classification Learner, only use a table that has a reduced number of features already, otherwise, it will take forever to load the data. (My computer took several minutes to load a table that was 162 x 39913). For computational reasons, you'll probably want max out at around 1000 (but this depends on your computer.)
- Matlab is probably looking for your "predictors" (features) to be in a table. If you need to convert an array to a table, check out: `array2table(brainvals)`.

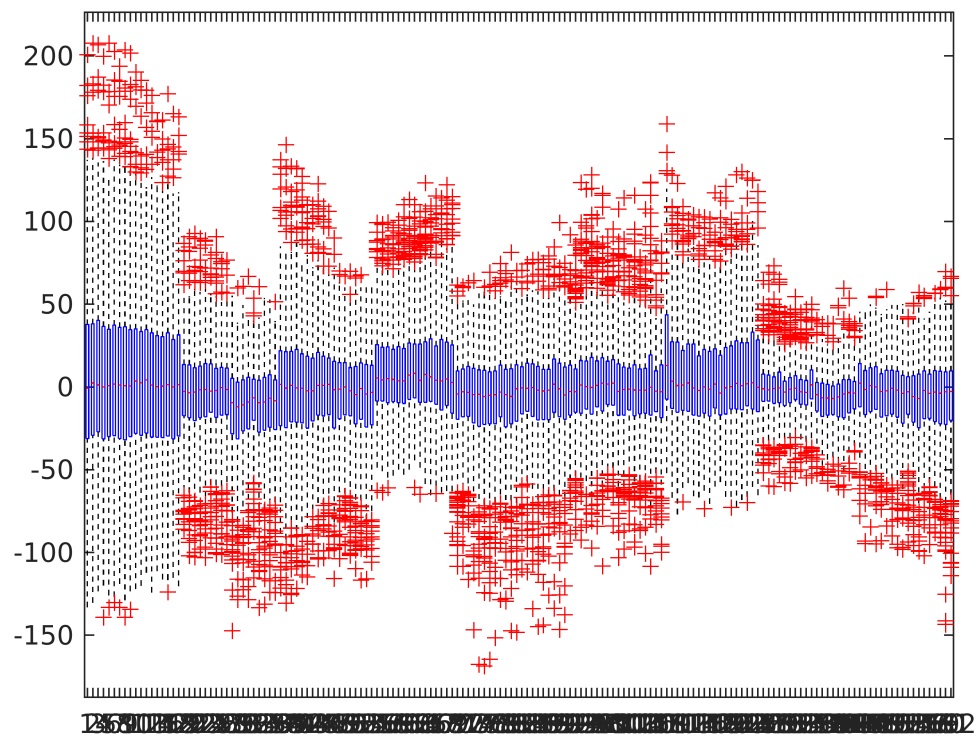
- If you are using Matlab 2019, you'll need to take one additional step if you want to use the classificationLearner app. You must put the “responses” and the “predictors” in the same table.
- You can do this by following this general format: `train_all = [table(train_labels) train_4regions ]`; In your version, you'll have a different variable than `train_4regions`.
- Or you can use something like this if you have an array:
- `train_table2d = array2table(train_fmri2d); %`
- `train_table2d.labels = train_labels;`
- Other tips from students (you can add a comment here).

```
n = 162;
region1 = fmri2d.*repmat(single(reshape(mask_region1,1,[])), [n 1]);
region1 = region1( :, any(region1))
```

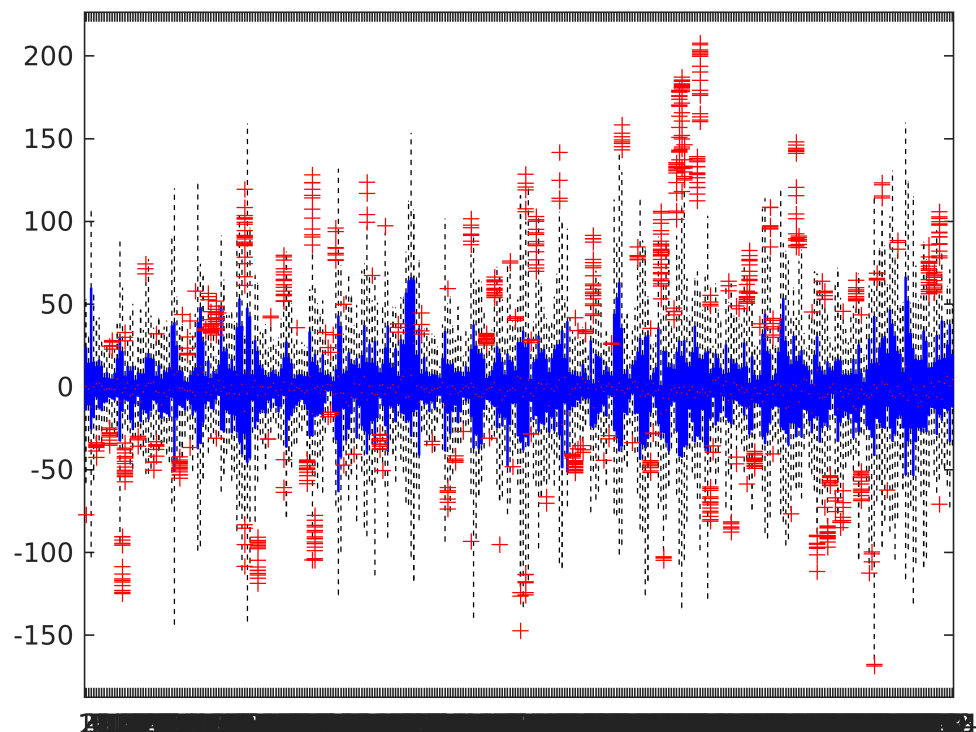
```
region1 = 162x334 single matrix
 44.8421    29.0900    88.2853   -23.1248   -0.8906    4.6762    25.0965   -6.1110 ...
 40.9222    27.1457   106.1361   -10.0830    7.1035    4.7111    23.0804    0.8875
 43.0023    30.2014    97.9869   -24.0412   -4.9025    0.7461    20.0644   -5.1141
 41.0824    25.2571    97.8377   -16.9994   -0.9084   -0.2190    20.0483   -5.1156
 40.1625    26.3128    99.6885   -12.9575    6.0856    4.8159    23.0322   -6.1171
 40.2426    24.3685    96.5393   -18.9157    2.0796    2.8509    19.0162   -8.1186
 49.3227    26.4242    82.3901   -20.8739    0.0737    1.8858    22.0001   -2.1201
 46.4028    27.4799    90.2409   -20.8321   -2.9323    0.9207    17.9841   -4.1216
 41.4829    22.5356    86.0917   -21.7903    0.0618    3.9556    20.9680   -5.1231
 42.2065    26.4295    70.0187   -19.3685    0.8591    1.1433    18.4219   -2.1741
  ⋮
```

```
% boxplot by sample
boxplot(region1')
```





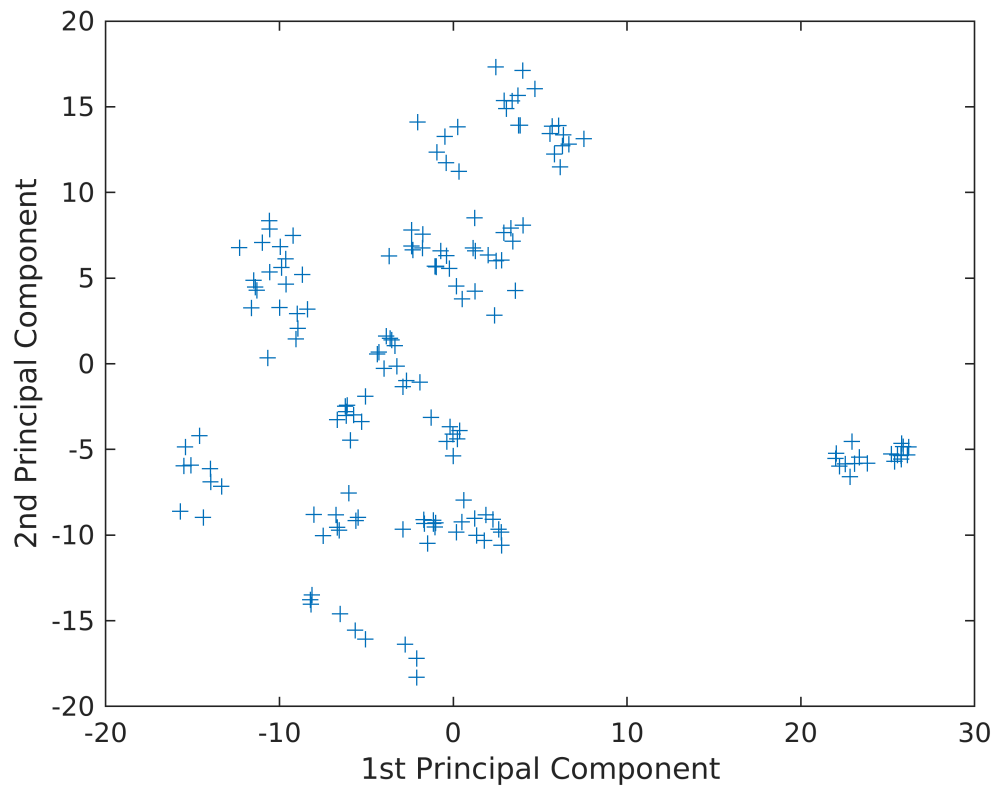
```
% boxplot by feature
boxplot(region1)
```



```

%corr(region1,region1);
[wcoeff,score,latent,~,explained] = pca(region1,'VariableWeights','variance');
figure()
plot(score(:,1),score(:,2),'+')
xlabel('1st Principal Component')
ylabel('2nd Principal Component')

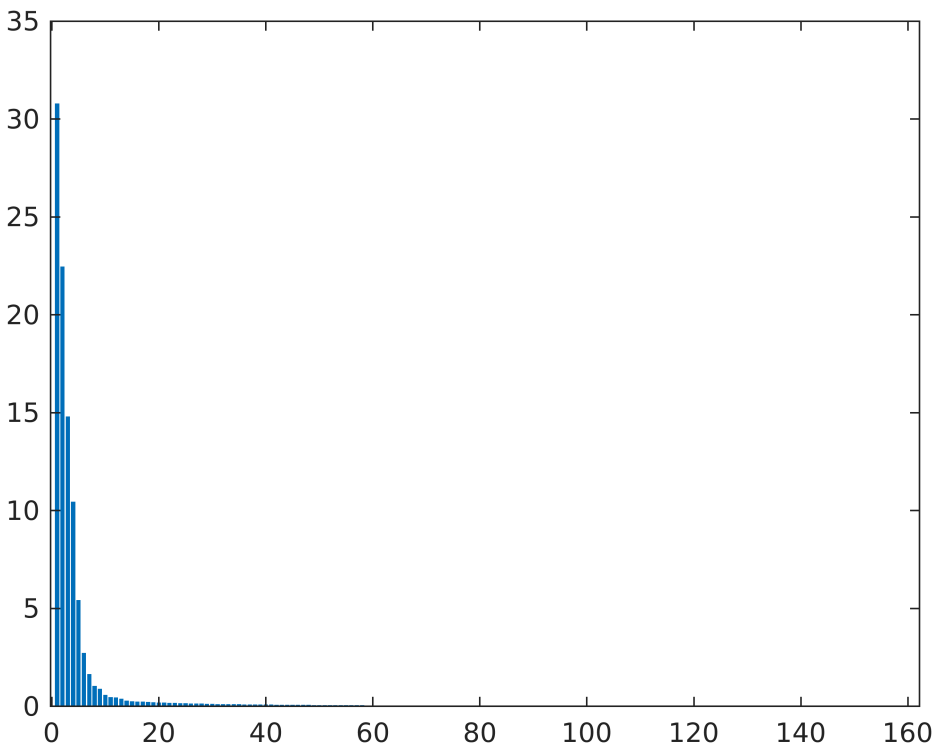
```



```

figure()
bar(explained)

```



```
region1 = fmri2d.*repmat(single(reshape(mask_region1,1,[])), [n 1]);
region1 = region1( :, any(region1));
region2 = fmri2d.*repmat(single(reshape(mask_region2,1,[])), [n 1]);
region2 = region2( :, any(region2));
region3 = fmri2d.*repmat(single(reshape(mask_region3,1,[])), [n 1]);
region3 = region3( :, any(region3));
region4 = fmri2d.*repmat(single(reshape(mask_region4,1,[])), [n 1]);
region4 = region4( :, any(region4));
```

```
[wcoeff1,score1,latent1,~,explained1] = pca(region1);
[wcoeff2,score2,latent2,~,explained2] = pca(region2);
```

Warning: Columns of X are linearly dependent to within machine precision.  
Using only the first 155 components to compute TSQUARED.

```
[wcoeff3,score3,latent3,~,explained3] = pca(region3);
[wcoeff4,score4,latent4,~,explained4] = pca(region4);
```

```
explained1 % first 7
```

```
explained1 = 161x1 single column vector
    36.2376
    23.9686
    15.8495
    10.7287
     4.2774
     1.8731
     1.1529
     0.7581
```

```
0.5554
0.4024
:
.
```

```
explained2 % first 8
```

```
explained2 = 161x1 single column vector
```

```
43.1001
19.9416
9.8243
7.8444
3.7835
2.1081
1.6599
1.2185
0.7526
0.5982
:
.
```

```
explained3 % first 9
```

```
explained3 = 161x1 single column vector
```

```
30.9446
19.2588
14.2063
5.7947
4.4339
3.8011
2.8119
1.3967
1.0155
0.9470
:
.
```

```
explained4 % first 8
```

```
explained4 = 31x1 single column vector
```

```
52.1313
15.8923
10.1483
7.3441
3.6160
2.2814
1.6386
1.3586
0.9103
0.7898
:
.
```

```
region1_10 = score1(:, 1:10)
```

```
region1_10 = 162x10 single matrix
```

```
882.2699 -249.0720 -17.4099 99.1775 28.8254 -62.8130 -19.3104 36.6531 ...
870.9958 -238.3671 0.1133 148.2624 38.4298 -62.0343 -25.2092 43.2560
884.9216 -229.5898 22.2092 141.0955 41.0693 -65.1118 -37.8074 31.4339
873.9255 -247.0448 -16.7304 142.4166 22.6208 -67.9866 -21.6513 26.2031
855.6765 -251.0692 -18.9221 132.2741 39.1652 -76.4831 -13.2179 27.8672
```

```

876.3199 -230.3638    3.4694 128.8664   34.3494 -67.4803 -25.3893   27.2125
866.0078 -256.0186    6.7014 100.1808   11.7942 -53.0865 -13.0122   30.2270
877.5620 -243.5495  -17.7944 103.2499   27.9399 -76.0061 -16.4831   19.0183
857.6501 -264.7130  -24.1600 118.0704   13.0859 -61.9725 -11.5007   23.3495
809.6169 -261.3054   67.0503   88.7841   25.6060 -41.0882  -6.1224   11.0350
⋮

```

```
region2_10 = score2(:, 1:10)
```

```
region2_10 = 162x10 single matrix
```

```

500.3787 -17.8039   53.0067 -63.0664   -6.9085 -30.1563   36.5739 -11.0019 ...
476.8148 -18.8309   38.2953 -64.6518  -12.3547  -7.2812   24.9917   1.2449
474.8785 -14.6288   43.4600 -61.0633  -21.0159 -10.7514   29.6061  -0.8047
491.4526 -25.5925   48.6052 -56.0446  -13.0624  -5.6371   27.1252  -6.8921
474.1767 -13.9762   48.3750 -63.4748  -13.0364  -5.8352   31.7320   7.9991
487.9001 -23.4758   43.8801 -59.3961  -15.3894  -26.1054   30.9382 -18.9955
490.4452 -26.3618   42.5140 -63.7428    3.0023 -41.1618   27.9324 -19.6379
486.4983 -28.6231   48.3556 -61.0351    3.5750 -55.3451   30.2215 -16.7835
504.4672 -30.7581   48.3434 -55.0466   -9.9613 -31.1667   20.1318 -13.8631
440.9283 -11.7494   18.2610 -34.5127   16.9274 -35.4605  -2.2187 -16.8447
⋮

```

```
region3_10 = score3(:, 1:10)
```

```
region3_10 = 162x10 single matrix
```

```

-602.9750    4.3032 -267.0129   -3.7728  -27.3058  -29.2039   46.4935   51.0825 ...
-606.9026   16.3646 -247.8228   39.7788  -34.4314  -10.7661   16.0021    4.3593
-597.9197   -5.4570 -248.6415   19.6119  -34.7483  -12.5068   44.4269   22.7052
-600.8988  -15.2894 -242.4704    3.1686  -50.0556  -32.8292   36.7889   22.3156
-610.7221    8.3663 -244.7362   10.9536  -52.8146  -28.6015   41.1013    8.9603
-605.3700  -13.4790 -243.7588   11.9757  -54.5816  -40.2087   64.9597    3.9702
-569.9920    2.7016 -245.7481    5.8671  -11.7410  -29.9437   26.8847   18.8572
-587.5948   -5.5146 -248.1978   27.1901  -25.3469  -60.6531   52.7184  -30.1970
-583.7650  -48.5553 -213.8648   24.3634  -29.9500  -65.3586   76.7826   -1.7185
-477.6545  -20.4015 -156.0078   -1.0636   58.3452  -51.4016  -21.7912   35.6059
⋮

```

```
region4_10 = score4(:, 1:10)
```

```
region4_10 = 162x10 single matrix
```

```

-200.7725  -88.9944   66.3656   30.2970    7.5523    2.0962   -5.4538  -5.7794 ...
-142.5894  -25.5757  128.2361   46.4931   -6.5584  -10.7303  -28.3200   9.1455
-169.1275  -21.1523  101.0748   87.2609  -56.4232   12.0617    2.9201   9.8378
-190.8640  -79.1459   77.3002   28.7072   -4.2127   -0.5531    6.3055  -0.1726
-185.5584  -79.0591   91.0915   35.9461   -4.6241   -1.3296    8.1930  -6.1525
-185.4122  -76.7435   78.5448   24.6691   -6.5634    6.8628   -1.1384  -0.2712
-197.2589  -59.4490   74.7602   43.4277   -8.1264   -1.1025  -22.3357   6.4737
-207.4717  -75.2307   78.9758   28.4065  -20.0961   -1.3167    1.7366 -11.4634
-212.6579  -75.7360   71.7959   24.7281   -7.0967   -1.8303   10.4450  -1.5655
-230.9032  -86.6549   53.2212   31.3413   -8.3701    1.6614  -15.4309  10.1136
⋮

```

```
chosen_features = [region1_10 region2_10 region3_10 region4_10]
```

```
chosen_features = 162x40 single matrix
```

```

882.2699 -249.0720  -17.4099   99.1775   28.8254  -62.8130  -19.3104   36.6531 ...
870.9958 -238.3671    0.1133  148.2624   38.4298  -62.0343  -25.2092   43.2560

```

```

884.9216 -229.5898 22.2092 141.0955 41.0693 -65.1118 -37.8074 31.4339
873.9255 -247.0448 -16.7304 142.4166 22.6208 -67.9866 -21.6513 26.2031
855.6765 -251.0692 -18.9221 132.2741 39.1652 -76.4831 -13.2179 27.8672
876.3199 -230.3638 3.4694 128.8664 34.3494 -67.4803 -25.3893 27.2125
866.0078 -256.0186 6.7014 100.1808 11.7942 -53.0865 -13.0122 30.2270
877.5620 -243.5495 -17.7944 103.2499 27.9399 -76.0061 -16.4831 19.0183
857.6501 -264.7130 -24.1600 118.0704 13.0859 -61.9725 -11.5007 23.3495
809.6169 -261.3054 67.0503 88.7841 25.6060 -41.0882 -6.1224 11.0350
⋮

```

```
covmat = cov(chosen_features)
```

```
covmat = 40x40 single matrix
```

```
Rows 27:36 | Columns 1:15
```

```
105 ×
```

```

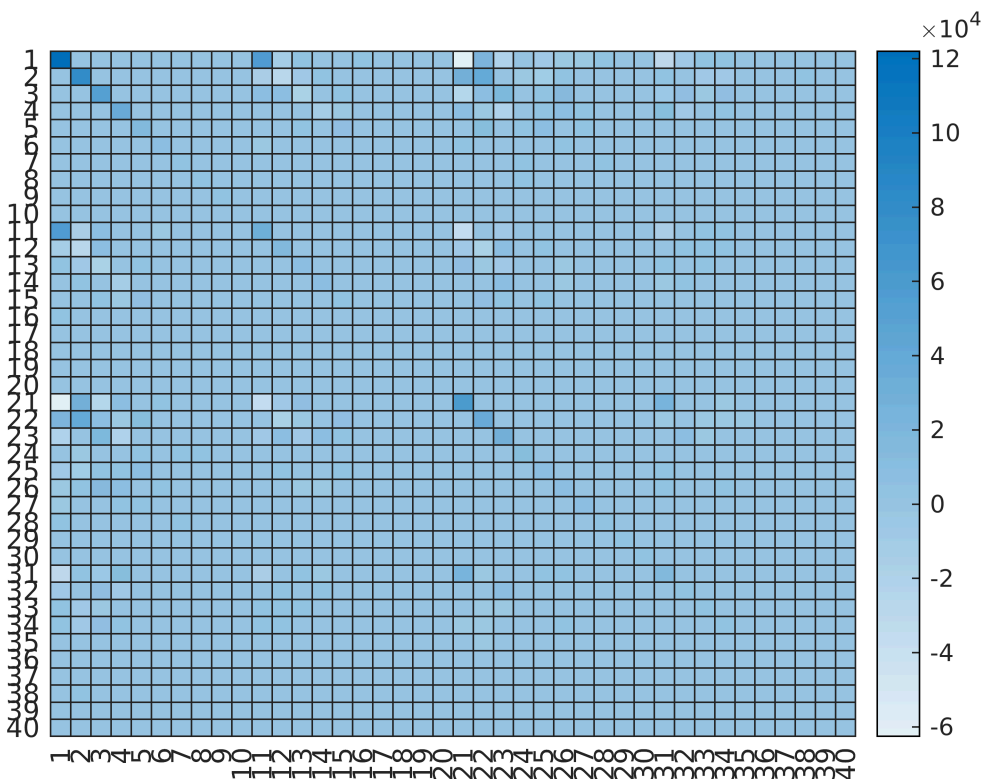
-0.0276 -0.0107 0.0130 0.0268 0.0363 0.0037 0.0034 0.0002
0.0228 -0.0080 -0.0071 -0.0048 -0.0031 0.0075 0.0177 -0.0047
0.0054 0.0144 0.0054 0.0026 -0.0097 -0.0159 0.0093 0.0036
0.0097 -0.0108 -0.0100 -0.0027 0.0098 0.0106 0.0036 0.0020
-0.3138 0.0361 -0.0614 0.1048 0.0101 0.0108 0.0049 0.0071
-0.0736 -0.0181 0.0517 -0.0632 -0.0173 0.0116 -0.0071 -0.0021
0.0437 -0.0746 -0.0543 0.0153 -0.0018 -0.0030 -0.0115 0.0004
0.0163 -0.0638 0.0482 0.0387 -0.0093 -0.0024 0.0020 -0.0006
-0.0061 -0.0192 -0.0226 0.0013 -0.0210 0.0054 0.0038 -0.0029
-0.0104 -0.0005 -0.0028 -0.0094 -0.0153 -0.0013 -0.0014 0.0054

```

```

figure();
heatmap(covmat)

```



After doing PCA and choosing the top 10 features from each region (after the first 10 features the explained variance < 1), I got a reduced feature set of 40 features.

Model	Accuracy
Fine Tree	88.9%
Medium Tree	88.9%
Coarse Tree	89.5%
Fine KNN	98.1% <-- pretty good training accuracy
Medium KNN	88.9%

## Brains on Trial [120 minutes]

Watch this video or read this paper and think about the prompts below.

- Brains on Trial 1 (60 minutes): <https://www.youtube.com/watch?v=o0eSqlAmKxU>
- Neuroimaging and lie detection paper: [https://drive.google.com/open?id=1V9YiLcl82VtDBnQLG\\_0-LT87tdW2D5tM](https://drive.google.com/open?id=1V9YiLcl82VtDBnQLG_0-LT87tdW2D5tM)

Other suggested resources if you want more of this:

- Brains on Trial: Part 2: [https://www.youtube.com/watch?v=\\_cBK\\_fgTZvk](https://www.youtube.com/watch?v=_cBK_fgTZvk)
- Law, Responsibility, and the Brain paper: <https://drive.google.com/file/d/1P-AoQPkPbp2EjC6u9LD-IXqgPjoRrChR/view?usp=sharing>

Prompt to start thinking about prior to class discussion:

How would you **go about determining** if a future, more accurate, version of brain-based lie detection technology should be permissible in the American justice system and in what situations?

What information would you want? (How realistic is it to obtain this information?) What groups would you need to consider? What measurements of tool effectiveness would you want to see (and under what conditions)? What social and societal implications should you consider?

You should write notes on the video and the prompt. You can either submit them (I would enjoy reading them), or write a note that you did this (what an appropriately themed use of the honor code!).

## Gati's Notes

- A jury is like a pooled decision maker
- I would want the data tested and fine-tuned to each individual
- **I would want all the jurors and judge to know the short-comings of the technology** (how do you make sure they listen and understand?)
- I would want masking techniques and adversarial patterns to be less effective